

Using Bayesian Networks to Predict Change Sets in an IT System

Sarah Nadi
snadi@uwaterloo.ca
CS886 - Fall 2010
University of Waterloo
Ontario, Canada

Abstract

When changing one component in an IT system, an analyst must be careful not to negatively impact any of the other system components. Detecting the complete change set (all those components that need to be changed to avoid any side effects) is therefore very important. Given historic information about how different components of a system change together, we build a Bayesian network that can be used to predict this change set. Given an initial component to be changed in the system, the posterior probability of the other components changing can be calculated using the network. In this paper, we experiment with different techniques for building the Bayesian network based on the data available, and report on their predictive power in terms of recall and precision.

1. Introduction

Change is inevitable in any Information Technology(IT) system. New features are added, different configurations are used, upgrades are introduced and new software and hardware are added. In a large system, such changes can cause other parts of the system to malfunction without the ability of the analyst to foresee this side effect. Accordingly, we need to ensure that we consider the impact of changes before they are implemented, and make sure we adjust the potentially affected components accordingly while implementing the change. This leads to the process of change set detection which is finding all the components of the system that need to be included in your change set. The change set is the set of components that need to be changed to avoid any side effects.

Many of the organizations with large and sophisticated IT systems employ a Configuration Management Database (CMDB) to help them manage these systems. The CMDB keeps track of all the components in a system, how they are related, as well as all the changes that have occurred

to them. Any component in the CMDB is called a Configuration Item (CI). In previous work [11], we mined the CMDB repository for historical co-changes, and used the mined correlations to predict change sets. That is, given an initial CI that is going to change, we predict what other CIs might need to be changed as well. We obtained promising results in terms of recall and precision.

In this paper, we wish to examine the same problem, but from a Bayesian perspective. Instead of looking at pairwise historical co-changes of CI, we would like to consider all the changed CIs, and learn the relationships between them. That is, given a set of observations about how different CIs change together, can we deduce the relationships between them that would allow us to predict future change sets. Accordingly, we explore the different ways a Bayesian network can be constructed from the data we have, and then test the predictions produced by querying this network to examine the obtained recall and precision. We hope that we can obtain better results using Bayesian networks since they are not limited to pairwise comparisons.

The rest of this paper is organized as follows. Section 2 provides the background information necessary for this work. It explains the nature of CMDBs, Bayesian networks, and the tools we used for analyzing Bayesian networks. Section 3 then presents the techniques we used to build our Bayesian network. This includes how we process the data, and build the different Bayesian networks accordingly. Section 4 details the setup we use for our experiments, and Section 5 provides the results obtained. Section 7 suggests future work that can be done to improve the obtained results. Section 8 concludes this paper.

2. Background

2.1. CMDBs and Change Sets

The repository of information we are using is the Configuration Management Database (CMDB) [12]. The CMDB is useful in Enterprise IT Management (EITM) since it provides information about the various critical components in a

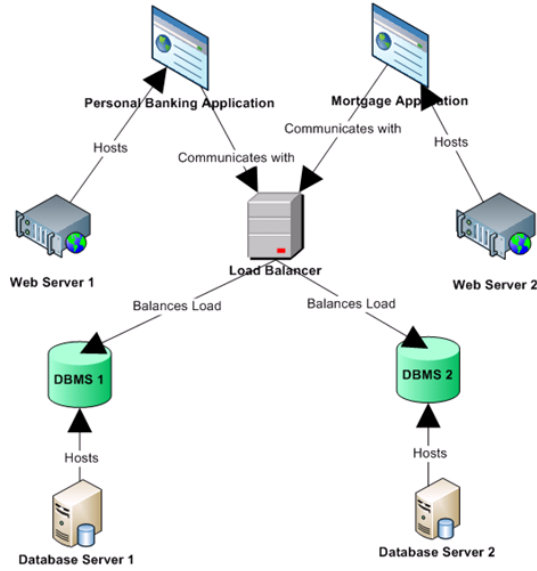


Figure 1. An IT System Stored in a CMDB. Each item shown is a CI, and CIs are tied through different relationship types

system including hardware, software, and services provided by the company. It records the configuration of these items, their change history, their incident history, as well as the relationships between them. Each item stored in the CMDB is referred to as a Configuration Item (CI). Figure 1 shows an example CMDB to illustrate the concepts of CIs and relationships.

A CMDB provides a basis for decision making processes such as Incident Management, Change Management, etc. In this paper, we focus on the process of Change Management, and in particular, on the problem of change set detection. A *change* is the addition, modification, or removal of anything that could affect IT services. A poorly planned change may lead to a fault in the system. Accordingly, when one wants to change a CI in the system, other CIs that might need to be changed as well must be correctly identified. We refer to the set of CIs that will need to be modified for the change to be complete as the *change set*.

2.2. Bayesian Networks

Bayesian Networks, first introduced by Pearl [13] are a type of probabilistic graphical models. Each node in the network represents a random variable, and nodes are connected through edges which represent probabilistic dependencies [2]. A Bayesian network is a Directed Acyclic Graph (DAG), and is essentially defined by two aspects: the structure of the graph and the Conditional Probability Tables (CPTs) probabilistically relating connected nodes. In our work, a node in the Bayesian network would be a CI which can take values “true”, “false” to denote the probability of it changing or not. The CPTs would essentially give

the probability of one CI changing given its parents’ values.

Given a set of data consisting of observations for our variable space, machine learning techniques can be used to estimate the network. If the structure of the network is known, and we have full observability (i.e we can observe values for each variable), then learning the CPTs can be achieved through Maximum-likelihood estimation (MLE) [2]. MLE tries to maximize the log likelihood of the the training data which can simply be done by calculating the frequency of occurrence of the values of each variable given its parents. If the structure is known, but we only have partial observability (i.e., missing data), then we can use a variation of MLE which is the Expectation Maximization algorithm [4]. If the structure is not known in the first place, then searching through the model space, and calculating network scores for each produced network is necessary. The simplest algorithm for this is a Greedy search algorithm [5].

While searching for the best network that fits the training data, scores are needed for the produced networks in order to know which one is the best one. There are several types of scores that have been proposed in the literature. Two of the commonly used scores are the Minimum Description Length (MDL) score [14] and the BDe score [8]. The MDL score is based on finding the length of data required to store a network. It tries to balance the complexity of the induced network with the degree of accuracy of which the network represents the frequencies in the training data. Therefore, the lower the MDL score of a network, the better it is. The BDe score is proportional to the posterior probability of each network structure being scored, given the data. Accordingly, we are searching for the network that maximizes this probability. For large data sets, both metrics score structures similarly [5].

2.3. Bayesian Network Tools

There are many machine learning and Bayesian analysis tools available. In this work, we use three of these tools which we briefly describe below.

2.3.1 WEKA

WEKA [6] is a Java toolbox for machine learning. It provides different algorithms for classification and clustering as well as other machine learning problems. It has Bayesian learning techniques for classification problems. Additionally, it provides explorations tools for Bayesian Networks, and allows the user to learn the structure and CPTs of a Bayesian Network. In terms of learning CPTs, one of the algorithms WEKA implements is the SimpleEstimator algorithm [16] which we use in our work. This simply estimates the probability values on edges based on the frequency values of each variable given its parents in the training set.

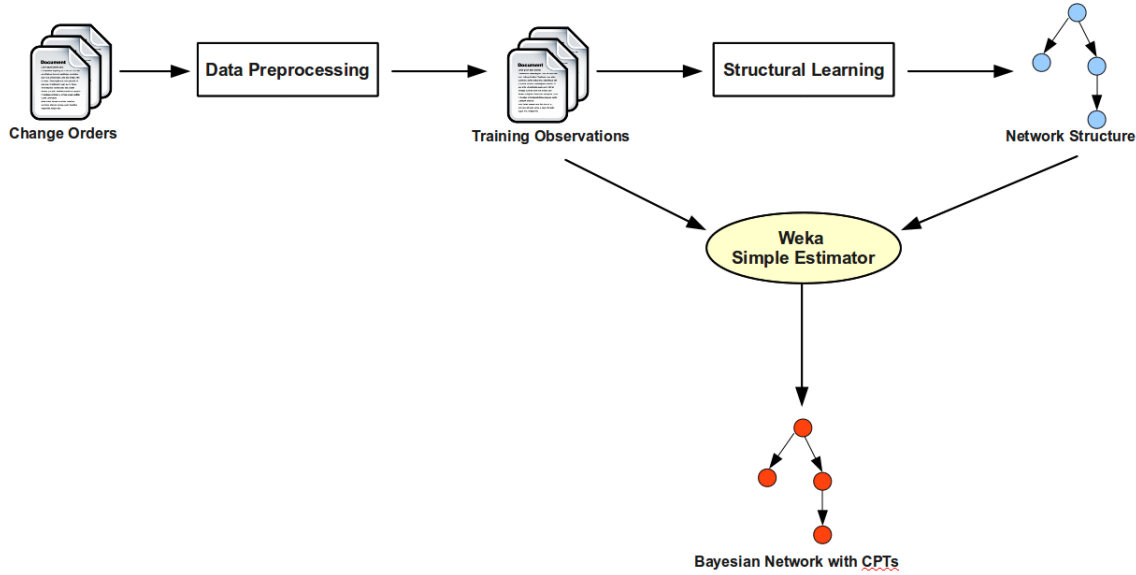


Figure 2. Building the Bayesian Network

2.3.2 Banjo

Banjo [1] is a tool written in Java which infers the structure of a Bayesian network given training data. It has two different search algorithms: Greedy and Simulated Annealing. Banjo uses the BDe metric to compute a network's score, and produces the highest scoring network found. We used the Greedy search algorithm in our experiments.

2.3.3 JavaBayes

JavaBayes [9] is a Java tool for calculating marginal probabilities and expectations in a Bayesian network. In a network, nodes can be set to be observed, and then the posterior probability of the remaining nodes can be calculated using Bayes theorem. This is the functionality we need here since the CI that will be changed will be considered as observed to be "true", and then we need to calculate the posterior of the remaining CIs given this observation.

3. Constructing the Bayesian Network

Before constructing the Bayesian Network, we first have to process and prepare the data we have. In order to construct a Bayesian network to use for predictions there are two steps involved. First, determining the structure of the actual network, and then estimating the Conditional Probability Tables (CPT). Section 3.1 first explains the data set available to us, and the data preprocessing involved before building the network. Section 3.2 then explains the different techniques we experimented with to build the network structure. Section 3.3 shows the last step to build the network which is estimating the CPTs. Figure 2 shows this

overall process.

3.1. Data Preprocessing

Data Size

The data set available to us in the sample CMDB we have is from three years (Jan 2006 - Dec. 2008), and has 7,999 distinct CIs, and 27,305 change orders. This amount of data was infeasible to work with as no tool could handle such a large amount of variables in a Bayesian network. For the purposes of this project, which is mainly to experiment with Bayesian techniques for change set prediction, it is sufficient to choose a small representative subset of the data. Accordingly, in order to be able to test things properly, we used observations from three months data from January 1, 2008 to March 31, 2008 to build the model. However, even in such a short period, there were already 2,841 distinct CIs appearing and 2,229 observations (i.e. change orders). Therefore, we needed to perform further data preprocessing.

First, we removed all CIs that have changed less than 12 times within this time frame (i.e. were associated with less than 12 different change orders in our training observations). There is no particular reason for choosing 12 as a cut off. It simply gave a feasible data set to deal with. This yielded 120 distinct CIs. Then, in order to slightly increase our variable space to include other related CIs, we found all the parent CIs related to these 120 CIs from the CMDB perspective (i.e. had a relationship edge in the CMDB). Adding the related parent CIs, we now had a set of 241 CIs. However, some of the added parent CIs may not be in the original set of 2,841 CIs. Accordingly, we just kept CIs that

appeared more than 12 times or were in the set of related CIs. This provided us our final data set of 170 CIs which we use throughout our models for fair comparison. Additionally, filtering out CIs meant filtering out some of the change orders that did not have any CIs satisfying our criteria (i.e. all the CIs that changed in them were outside our 170 CI set). This led to us having 1,305 change orders instead of 2,229 which was more manageable set. Accordingly, our training data set consisted of 170 variables (CIs), and 1,305 observations (change sets).

Data Format

In the CMDB, a Change Order has several fields including the requester, the assignee, the start date of the change, the description and the change set field (called the 'Configuration items' fields). Of the fields in a change order, we use only the change set field. The advantage of only using this one field is that change sets are easy to extract and easy to understand. In terms of an observation, CIs were either marked as "true" or "false" indicating whether they have changed or not. Each observation initially consisted of all the CIs in our data set marked as "false". For each change order in the training set, we would mark each CI appearing in the "Configuration Items" field as true, while all those not appearing were left as "false". Therefore, the set of training data consisted of an observation entry for each change order with the CIs marked as "true" or "false" accordingly.

3.2. Network Structure

There were essentially three different ways in which we could estimate the structure of the network. We, thus, built a Bayesian network using each technique to be able to compare their performance.

3.2.1 Method 1 (Banjo)

For the first method, we used Banjo to learn the structure of the Bayesian network from the training data. We use the Greedy searching algorithm implemented there. Banjo produced five different top-scoring networks (having the same score), and we simply chose one of them as the network to be used.

3.2.2 Method 2 (WEKA Classification)

Although our research problem as currently defined is not really a classification problem, we wanted to examine how the network built by the BayesNet classifier in WEKA will perform. Accordingly, we estimated the structure of the network using the K2 structure learning algorithm implemented in WEKA which is also a greedy search based learning algorithm. The default maximum number of parents allowed for any node is one (a Naive Bayesian Network).

However, we used a maximum of five to allow the nodes to be related since there is really no class label in our data set. We could not experiment with other structural learning techniques in WEKA because they do not allow multiple parents.

3.2.3 Method 3 (CMDB Relationships)

For the third method, we used the relations existing in the CMDB to infer the structure of the network. We used two variations for this. The first was to follow the direction of the relationship edges in the CMDB. That is, if there is a relationship edge in the CMDB from A to B, we will place an edge in the Bayesian network from A to B. The second was to reverse the direction of the relationship edges in the CMDB. That is, if there is a relationship edge from A to B in the CMDB, we will place an edge in the Bayesian network from B to A. There was one problem, however, with the generated networks. Two of the CIs had more than 15 parents to them which made their CPTs more intractable to compute. For those two CIs, we simply removed all their parents to make the computation tractable.

3.3. Estimating the CPTs

After the structure was determined, we used the SimpleEstimator algorithm [16] built in WEKA to calculate the CPTs for all networks. After setting the data set to be the observations in our training set (an ARFF file), WEKA learned the CPTs for each of the networks above. This complete network was then saved as an BIF XML file to be used by JavaBayes for general inference as will be explained in the next section.

Additionally, we experimented with the missing data feature in WEKA. WEKA allows the user to specify missing values for any variable by simply putting a '?' instead of its value. Therefore, instead of marking the CIs that did not appear in a change order as "false", we put a "?" in their place. This seemed closer to practice, because in reality, we are not sure whether this CI actually changed or not. When using Banjo or the CMDB relationships to estimate the structure of the network, the missing values were only inserted when learning the CPTs. When using WEKA, the missing values were inserted both to learn the network, and to estimate the CPTs.

4. Experiment Setup

This section explains how we tested the predictions of each of the networks in a way that simulates how an analyst would use the tool in reality. We first explain how we simulate this behavior, and then explain how we evaluate the predicted CIs. Algorithm 1 summarizes these steps taken to predict a change order.

Algorithm 1 Generating Predictions using the Bayesian Network

Input: *changeOrder*
Input: *BayesianNetwork*
Input: *networkCIs*
Input: *threshold*
Set *initialCI* = first CI in *changeOrder*
Set *occurredset* = *changeOrder* - *initialCI*
Initialize *observedSet* = {*initialCI*}
Initialize *predictedSet* = {}
repeat
 Initialize *newPredictions* = {}
 if *initialCI* not in *networkCIs* **then**
 return
 end if
 for *node* in *BayesianNetwork* **do**
 posterior = query *node* in *BayesianNetwork*
 if *posterior* > *threshold* **and** *node* not in *predictedset* **then**
 Add *node* to *newPredictions*
 end if
 end for
 for *prediction* in *newPredictions* **do**
 if *prediction* in *occurredSet* **then**
 Add *prediction* to *observedSet*
 end if
 end for
 Add *newPredictions* to *predictedSet*
until *newPredictions* is {}

$$recall = \frac{predictedSet \cap occurredSet}{occurredSet}$$
$$precision = \frac{predictedSet \cap occurredSet}{predictedSet}$$

4.1. Change Set Detection Process Simulation

The change set detection process provided is an iterative, collaborative process between the tool and the analyst. When the tool suggests CIs, the analyst can accept or reject these CIs (i.e. add them to the change set or not). Based on the CIs added to the change set, the analyst can ask the tool for more suggestions. This is along the lines of “Now, that I am also going to change these CIs, what else do I need to change?”. This process is explained in more details by Nadi et al. [11].

To simulate this process for our experiments, we do the following. At this point, we have the Bayesian network ready, and we would like to perform Inference. More formally, given that a CI will change (our observation), we want to update our beliefs that the other CIs in the network might change as well. Unfortunately, neither of the two tools previously used provide a general inference engine.

We, therefore, use JavaBayes in this step since it accepts the same XML BIFF format produced by WEKA. We use a one month testing set where we try to predict all the change orders in April 2008. There was a total of 883 change orders in that month which we use in our testing set. No preprocessing was performed on the testing data.

For each change order, we would take the first CI recorded as the initial CI to change, then we would set that as an observed node, and update the beliefs of all the nodes (CIs) in the network using JavaBayes. We are basically checking the posterior probability of each CI being “true” given our observation. We used different cutoff thresholds for the posterior probabilities we would consider. We would then loop on all the updated CIs, and add those that match our threshold criteria to the predicted change set. For example a threshold of 0.2 means that only nodes that have a posterior probability greater than 0.2 of having the value “true” will be returned. To simulate a real life scenario, we then checked which of these predicted CIs actually lies in the target change set we are trying to predict. This is similar to an analyst accepting CIs into their change set. These common CIs would then also be marked as observed so that we can predict what else will need to change given that these CIs are also changing. Again, we would calculate the posterior probability given the new observations, and continue doing so until there are no more new CIs to predict. All the CIs that match the threshold criteria (whether accepted by the analyst or not) are part of the predicted set. These steps (shown in Algorithm 1) would be repeated for each change order. Note that if the initial CI is not in the list of CIs we used for training, then the predicted set will be empty since we have not seen this CI before.

4.2. Evaluation Techniques

We need a way to evaluate the predicted CIs. The recall and precision measures from Information Retrieval are appropriate for this type of evaluation. Recall measures the proportion of correct CIs retrieved by the system, while precision measures the proportion of suggested CIs that are correct [15].

Similar to Hassan et. al [7], we define the *Predicted Set* (P) as the set of all predicted. We define the *Occurred Set* (O) as the CIs remaining in the change set after excluding the Initial CI provided by the analyst (i.e Change Set - Initial CI). The intersection of the predicted set and the occurred set, called *PO*, is the common CIs in both sets. For each constructed change set, we then calculate the recall and precision values for the predictions according to the following definitions [7]:

$$Recall = \frac{|PO|}{|O|} \quad (1)$$




Figure 3. F-measure of the different Bayesian networks produced using the CMDB relationships

$$Precision = \frac{|PO|}{|P|} \quad (2)$$

If no CIs are predicted (i.e., P and thus PO are empty), precision is defined as 1 since there cannot exist any incorrect predictions in an empty set. On the other hand, if the size of the change set is 1, and thus the size of the occurred set is 0, recall is defined as 1 since there are no CIs to predict [7].

In order to have a single measure that indicates the effectiveness of our predictions, we use the F-measure which is based on van Rijsbergen’s effectiveness measure which combines recall and precision [15]. The F-measure is calculated according to Equation 3 which gives equal weighting to recall and precision. The ideal F-measure is 1 where both recall and precision are 1.

$$F = 2 * \frac{precision * recall}{precision + recall} \quad (3)$$

5. Results

For all the networks generated from the different methods explained in Section 3.2, and following the procedure described in Section 4, we calculated the average recall and precision for all the change orders in the test set. Based on the average recall and precision values, we calculated the F-measure. Table 1 summarizes the recall and precision obtained for all models rounded to three decimal places. For better visualization, we plot the F-measure from all the networks based on CMDB structure together, and the F-measure from those based on Banjo and WEKA together.

Figure 3 shows the F-measure obtained from the four different networks built based on CMDB relationships for structural information. The first observation is that reversing the relationship edges did not produce any difference in the results. The curves for the CMDB relationships and the reversed CMDB relationships are completely overlapping.

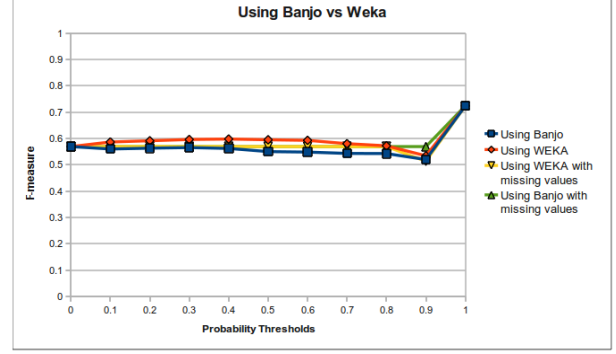


Figure 4. F-measure from network structures estimated by WEKA versus Banjo

The same things applies to the curves of CMDB relationships with missing data, and reversed CMDB relationships with missing data. The second observation here is that denoting CIs that did not appear in the change order as missing data rather than giving them the value “false” produced slightly better results. However, the main conclusion from these results is that given each change order, all of the predicted CIs in the network already had a posterior probability of greater than 0.9. That is why the F-measure is the same for all thresholds (excluding 0), and only changes when the threshold becomes 0.9. This seems rather strange, and despite the relatively high F-measure, it does not seem that the structure of these networks is reflective of their relationships in terms of change.

Figure 4 compares the F-measure obtained from the graphs whose structure was built using Banjo versus those built using WEKA. It also shows the effects including missing data in both cases. Surprisingly, it seems that the network built using WEKA has better performance (although only slightly better). However, marking CIs that have not been changed as unobserved (missing) instead of “false” slightly deteriorated the predictive performance. Overall, the performance of the two networks are very similar which comes as a bit of a surprise since WEKA is really a classification engine and assumes that the last variable in the data is the target class which is not the case in our data.

When comparing networks built from CMDB structure, and those built by WEKA or Banjo, we still notice that the performance in terms of the F-measure is very close. However, WEKA and Banjo had slightly better recall as shown in Table 1. Additionally, while running our tests, it took much longer to query the graphs built by WEKA and Banjo than those built from the CMDB structure. This implies that the former graphs have more interconnections between the nodes while the latter are more sparse.

We also note when the threshold is 1.0 (i.e no predictions are being made), the recall is not 0 although the precision is 1. This is because there are many change sets of size 1

Threshold	Banjo		Banjo Missing		WEKA		WEKA Missing		CMDB		CMDB Missing		Reversed CMDB		Reversed CMDB Missing	
	Recall	Precision	Recall	Precision	Recall	Precision	Recall	Precision	Recall	Precision	Recall	Precision	Recall	Precision	Recall	Precision
0	0.715	0.473	0.715	0.473	0.715	0.473	0.715	0.473	0.715	0.473	0.715	0.473	0.715	0.473	0.715	0.473
0.1	0.649	0.493	0.715	0.473	0.707	0.501	0.715	0.473	0.569	0.470	0.715	0.473	0.583	0.471	0.715	0.473
0.2	0.643	0.501	0.715	0.473	0.704	0.509	0.715	0.473	0.568	0.470	0.715	0.473	0.580	0.470	0.715	0.473
0.3	0.634	0.510	0.715	0.473	0.700	0.519	0.715	0.473	0.568	0.469	0.715	0.473	0.569	0.469	0.715	0.473
0.4	0.624	0.511	0.715	0.473	0.696	0.524	0.715	0.473	0.568	0.469	0.715	0.473	0.568	0.469	0.715	0.473
0.5	0.610	0.500	0.715	0.473	0.682	0.528	0.715	0.473	0.568	0.469	0.715	0.473	0.568	0.469	0.715	0.473
0.6	0.607	0.500	0.715	0.473	0.674	0.529	0.715	0.473	0.567	0.469	0.715	0.473	0.568	0.469	0.715	0.473
0.7	0.595	0.499	0.715	0.473	0.646	0.526	0.715	0.473	0.567	0.469	0.715	0.473	0.568	0.469	0.715	0.473
0.8	0.595	0.499	0.715	0.473	0.627	0.526	0.715	0.473	0.567	0.469	0.715	0.473	0.567	0.469	0.714	0.473
0.9	0.574	0.475	0.715	0.473	0.588	0.489	0.567	0.469	0.567	0.469	0.715	0.473	0.567	0.469	0.712	0.473
1	0.567	1.000	0.567	1.000	0.567	1.000	0.567	1.000	0.567	1.000	0.567	1.000	0.567	1.000	0.567	1.000

Table 1. Summary of Results

which makes the occurred set of size 0, and thus recall being automatically 1. Please refer to Nadi et. al [11] for more details. Additionally, all networks start at the same recall and precision values when the threshold is 0 since simply all 170 CIs are predicted. This is because in the Bayesian network, the probability of any CI is never really zero. It may be very close to zero, but will never be completely unlikely.

Given that there will be many CIs in the testing data that have not been seen before in the training, we thought of limiting the test data to only those nodes that appeared in the training, and adjusting the observations in the test set accordingly. In that case, for almost all models, we would either see perfect recall and almost zero precision or perfect precision and the recall obtained when no predictions are made. This indicates that limiting the test data to the same set of CIs used in training leads to some sort of overfitting. These results are not shown here due to space limitations.

6. Related Work

Mirarab et al. [10] investigate the same problem as our work. They build three different Bayesian Networks, one that is based on package and class dependency information (static relationships), one which is dependent on historical co-changes, and one which uses both. For the first graph, the initial structure is essentially “given” according to the static dependencies in the code, and then the CPTs are learnt using the importance sampling algorithm proposed by Changhe and Marek [17]. The way static dependencies are defined in their case is specific to Java. The third one is essentially the first graph, but updated using the historic change information according to the Expectation Maximization (EM) algorithm [4]. The second was solely based on historic information where the network is build using a greedy structure learning algorithm [5]. The main difference between their work and ours is that their work is on the level of source code while ours is on a system level.

Zhou et al. [18] try to answer a slightly different problem. They do not only look at the probability of other elements changing given a specific element, they also add features such as authors, change significance levels etc. and try to predict if two elements are co-changes or not accordingly. Thus, their problem is more of a classification prob-

lem where given two elements, and some observed features they try to determine the class as a co-change or not. They use the K2 algorithm proposed by Cooper et. al [3] to estimate the structure of the Bayesian network, and use the SimpleEstimator algorithm built in WEKA [16]. Although the problems being solved are different, we used similar almost the same techniques they used to build the network.

7. Future Work

There is still a lot of work that needs to be done in this area. First, it would be interesting to experiment with undirected graphs. For our problem, when two CIs, A and B, occur in the same change order, it is hard to determine whether A caused B to change or B caused A to change. Therefore, using an undirected model, such as a Markov network could be interesting in this setting. Unfortunately, due to time constraints, we could not experiment with that. Another possible expansion is to use the CMDB relationships more carefully when determining the structure of the network. Currently, we either add edges to the Bayesian network following the direction of those in the CMDB or in the reverse direction. It would be interesting to determine which set of relationships semantically make sense to have the causality relationship in their same direction, and which make sense to have it in the opposite direction. That way, the structure of the network produced could perhaps be more accurate.

Another addition would be to make the model more “intelligent”. That is, instead of not predicting anything at all for CIs that have not been previously seen, it can find a similar CI it has information about, and predict the change set based on it. The challenge here is how to define “similar”. This can simply be the same type of CI (i.e. database server, application, router etc.) with the assumption that CIs of the same type will behave similarly or it can be CIs with the same relationships. This would require further investigation as to which criteria is the best for describing similarity in this context.

8. Conclusion

In this work, we experimented with Bayesian networks in the context of change set predictions for IT systems. Given the Configuration Items (CIs) the analyst wants to change in the system, we can predict what other CIs might need to be changed as well. We built different Bayesian networks based on different information such as CMDB relationships or using structural learning algorithms such as the Greedy algorithm. Overall, all the networks built gave us, more or less, the same results (an F-measure of about 0.6). We believe that this is a good start for this exploratory work which indicates that Bayesian networks can be employed in change set detection. However, more work needs to be done in terms of making better predictions for unseen CIs, and in terms of exploring other types of probabilistic graphs that might be used in this context. We hope to expand this work to explore more variations of Bayesian learning in order to get better change set predictions.

References

- [1] Banjo. "<http://www.cs.duke.edu/~amink/software/banjo/>".
- [2] I. Ben-Gal. Bayesian networks.
- [3] G. Cooper and E. Herskovits. A Bayesian method for the induction of probabilistic networks from data. *Machine learning*, 9(4):309–347, 1992.
- [4] A. Dempster, N. Laird, D. Rubin, et al. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society. Series B (Methodological)*, 39(1):1–38, 1977.
- [5] N. Friedman and M. Goldszmidt. Learning Bayesian Networks with Local Structure. 1996.
- [6] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. Witten. The WEKA data mining software: An update. *ACM SIGKDD Explorations Newsletter*, 11(1):10–18, 2009.
- [7] A. E. Hassan and R. C. Holt. Predicting change propagation in software systems. In *ICSM '04: Proceedings of the 20th IEEE International Conference on Software Maintenance*, pages 284–293, Washington, DC, USA, 2004. IEEE Computer Society.
- [8] D. Heckerman, D. Geiger, and D. Chickering. Learning Bayesian networks: The combination of knowledge and statistical data. *Machine learning*, 20(3):197–243, 1995.
- [9] Javabayes version 0.346 [online]. "<http://www-2.cs.cmu.edu/~javabayes/Home/>".
- [10] S. Mirarab, A. Hassouna, and L. Tahvildari. Using bayesian belief networks to predict change propagation in software systems. pages 177–188, jun. 2007.
- [11] S. Nadi, R. Holt, and S. Mankovskii. Does the past say it all? Using history to predict change sets in a CMDB. In *CSMR'10: Proceedings of the 14th European Conference on Software Maintenance and Reengineering*, Madrid, Spain, 2010.
- [12] Office of government commerce (ogc), ed.: Service support. *IT Infrastructure Library (ITIL)*, 2000.
- [13] J. Pearl. *Probabilistic reasoning in intelligent systems: networks of plausible inference*. Morgan Kaufmann, 1988.
- [14] J. Rissanen. Stochastic complexity and statistical inference. *Analysis and Optimization of Systems*, pages 391–407, 1986.
- [15] C. van Rijsbergen. Information retrieval, 1979.
- [16] I. Witten and E. Frank. *Data Mining: Practical machine learning tools and techniques*. Morgan Kaufmann Pub, 2005.
- [17] C. Yuan and M. Druzdzel. An importance sampling algorithm based on evidence pre-propagation. In *Proceedings of the 19th Annual Conference on Uncertainty on Artificial Intelligence*, pages 624–631. Cite-seer, 2003.
- [18] Y. Zhou, M. Würsch, E. Giger, H. C. Gall, and J. L? A bayesian network based approach for change coupling prediction. *Reverse Engineering, Working Conference on*, 0:27–36, 2008.